

Кузьмин Роман ИУ5-64Б Вариант №12

```
[*]: import numpy as np
import pandas as pd
from typing import Dict, Tuple
from scipy import stats
from IPython.display import Image
from sklearn.datasets import load_iris, load_boston
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, balanced_accuracy_score
from sklearn.metrics import precision_score, recall_score, f1_score, classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import mean_absolute_error, mean_squared_error, mean_squared_log_error, median_absolute_error, r2_score
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.preprocessing import MinMaxScaler
from sklearn.datasets import make_blobs, make_circles
from sklearn.model_selection import cross_val_score, cross_validate
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC, NuSVC, LinearSVC, OneClassSVM, SVR, NuSVR, LinearSVR
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import RandomizedSearchCV
from sklearn.ensemble import AdaBoostClassifier
from sklearn import svm
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style="ticks")
```

```
[2]: data = pd.read_csv('marvel-wikia-data.csv', sep=',')
```

```
[3]: data.head()
```

	page_id	name	urlslug	ID	ALIGN	EYE	HAIR	SEX	GSM	ALIVE	APPEARANCES	FIRST APPEARANCE	Year
0	1678	Spider-Man (Peter Parker)	VSpider-Man_(Peter_Parker)	Secret Identity	Good Characters	Hazel Eyes	Brown Hair	Male	NaN	Living Characters	4043.0	Aug-62	1962.0
1	7139	Captain America (Steven Rogers)	VCaptain_America_(Steven_Rogers)	Public Identity	Good Characters	Blue Eyes	White Hair	Male	NaN	Living Characters	3360.0	Mar-41	1941.0
2	64786	Wolverine (James "Logan" Howlett)	\Wolverine_(James_%22Logan%22_Howlett)	Public Identity	Neutral Characters	Blue Eyes	Black Hair	Male	NaN	Living Characters	3061.0	Oct-74	1974.0
3	1868	Iron Man (Anthony "Tony" Stark)	VIron_Man_(Anthony_%22Tony%22_Stark)	Public Identity	Good Characters	Blue Eyes	Black Hair	Male	NaN	Living Characters	2961.0	Mar-63	1963.0
4	2460	Thor (Thor Odinson)	VThor_(Thor_Odinson)	No Dual Identity	Good Characters	Blue Eyes	Blond Hair	Male	NaN	Living Characters	2258.0	Nov-50	1950.0

```
[4]: data.isnull().sum()
```

```
[4]: page_id      0
name        0
urlslug      0
ID         3770
ALIGN       2812
EYE        9767
HAIR       4264
SEX         854
GSM        16286
ALIVE        3
APPEARANCES 1096
FIRST_APPEARANCE 815
Year        815
dtype: int64
```

```
[5]: data.shape
```

```
[5]: (16376, 13)
```

```
[6]: data.pop('GSM')
```

```
[6]: 0      NaN
1      NaN
2      NaN
3      NaN
4      NaN
...
16371   NaN
16372   NaN
16373   NaN
16374   NaN
16375   NaN
Name: GSM, Length: 16376, dtype: object
```

```
[7]: data.shape
```

```
[7]: (16376, 12)
```

```
[8]: data = data.dropna(axis=0, how='any')
data.shape
```

```
[8]: (4402, 12)
```

```
[9]: data.head()
```

	page_id	name	urlslug	ID	ALIGN	EYE	HAIR	SEX	ALIVE	APPEARANCES	FIRST APPEARANCE	Year
0	1678	Spider-Man (Peter Parker)	VSpider-Man_(Peter_Parker)	Secret Identity	Good Characters	Hazel Eyes	Brown Hair	Male Characters	Living Characters	4043.0	Aug-62	1962.0
1	7139	Captain America (Steven Rogers)	VCaptain_America_(Steven_Rogers)	Public Identity	Good Characters	Blue Eyes	White Hair	Male Characters	Living Characters	3360.0	Mar-41	1941.0
2	64786	Wolverine (James "Logan" Howlett)	Wolverine_(James_%22Logan%22_Howlett)	Public Identity	Neutral Characters	Blue Eyes	Black Hair	Male Characters	Living Characters	3061.0	Oct-74	1974.0
3	1868	Iron Man (Anthony "Tony" Stark)	VIron_Man_(Anthony_%22Tony%22_Stark)	Public Identity	Good Characters	Blue Eyes	Black Hair	Male Characters	Living Characters	2961.0	Mar-63	1963.0
4	2460	Thor (Thor Odinson)	VThor_(Thor_Odinson)	No Dual Identity	Good Characters	Blue Eyes	Blond Hair	Male Characters	Living Characters	2258.0	Nov-50	1950.0

Кодируем категориальные признаки

```
[10]: data.dtypes
```

```
[10]: page_id          int64
      name            object
      urlslug         object
      ID              object
      ALIGN           object
      EYE             object
      HAIR            object
      SEX              object
      ALIVE            object
      APPEARANCES     float64
      FIRST_APPEARANCE object
      Year             float64
      dtype: object
```

```
[11]: from sklearn.preprocessing import LabelEncoder, OneHotEncoder
le = LabelEncoder()
df_int = le.fit_transform(data['name'])
data['name'] = df_int
df_int = le.fit_transform(data['urlslug'])
data['urlslug'] = df_int
df_int = le.fit_transform(data['ID'])
data['ID'] = df_int
df_int = le.fit_transform(data['ALIGN'])
data['ALIGN'] = df_int
df_int = le.fit_transform(data['EYE'])
data['EYE'] = df_int
df_int = le.fit_transform(data['HAIR'])
data['HAIR'] = df_int
df_int = le.fit_transform(data['SEX'])
data['SEX'] = df_int
df_int = le.fit_transform(data['ALIVE'])
data['ALIVE'] = df_int
df_int = le.fit_transform(data['FIRST_APPEARANCE'])
data['FIRST_APPEARANCE'] = df_int
data.head()
```

```
[11]:   page_id  name  urlslug  ID  ALIGN  EYE  HAIR  SEX  ALIVE  APPEARANCES  FIRST_APPEARANCE  Year
0    1678  3738  3738  3    1     8     5     3    1    4043.0          80  1962.0
1    7139   624   624  2    1     3    20     3    1    3360.0          425  1941.0
2   64786  4302  4302  2    2     3     2     3    1    3061.0          622  1974.0
3   1868  1785  1785  2    1     3     2     3    1    2961.0          434  1963.0
4    2460  3942  3942  1    1     3     3     3    1    2258.0          548  1950.0
```

Масштабируем числовые данные

```
[12]: sc1 = MinMaxScaler()
data['page_id'] = sc1.fit_transform(data[['page_id']])
data['APPEARANCES'] = sc1.fit_transform(data[['APPEARANCES']])
data['Year'] = sc1.fit_transform(data[['Year']])
data.head()
```

```
[12]:   page_id  name  urlslug  ID  ALIGN  EYE  HAIR  SEX  ALIVE  APPEARANCES  FIRST_APPEARANCE  Year
0  0.000866  3738  3738  3    1     8     5     3    1    1.000000          80  0.310811
1  0.008108   624   624  2    1     3    20     3    1    0.831024          425  0.027027
2  0.084554  4302  4302  2    2     3     2     3    1    0.757051          622  0.472973
3  0.001118  1785  1785  2    1     3     2     3    1    0.732311          434  0.324324
4  0.001903  3942  3942  1    1     3     3     3    1    0.558387          548  0.148649
```

```
[13]: data['ALIVE'].unique()
[13]: array([1, 0])
```

Разделение на выборки

```
[14]: target = data['ALIVE']
data_X_train, data_X_test, data_y_train, data_y_test = train_test_split(
    data, target, test_size=0.2, random_state=1)

[15]: data_X_train.shape, data_y_train.shape

[15]: ((3521, 12), (3521,))

[16]: data_X_test.shape, data_y_test.shape

[16]: ((881, 12), (881,))

[17]: np.unique(target)

[17]: array([0, 1])
```

Логистическая регрессия

```
[ ]: model = LogisticRegression()
model.fit(data_X_train, data_y_train)

[21]: data_y_pred = model.predict(data_X_test)
accuracy_score(data_y_test, data_y_pred)

[21]: 1.0

[22]: f1_score(data_y_test, data_y_pred, average='micro')

[22]: 1.0
```

Градиентный бустинг


```
[22]: xgb_model = xgb.XGBClassifier()
xgb_params = {'learning_rate': [0.01, 0.05, 0.1, 0.12, 0.15, 0.2, 0.3, 0.5, 0.7],
               'max_depth': [5, 6, 7, 8],
               'min_child_weight': [2, 3, 5, 11],
               'n_estimators': [1, 2, 5, 10, 12, 15, 20]}

xgb_search = GridSearchCV(estimator=xgb_model, param_grid=xgb_params, cv=5, n_jobs=4, scoring='accuracy')
xgb_search.fit(train_X, train_y)
xgb_search.best_params_, xgb_search.best_score_

c:\Users\danih\Documents\python\myvenv\lib\site-packages\sklearn\model_selection\_split.py:676: UserWarning: The least populated class in y has only 2 members, which is less than n_splits=5.
warnings.warn(
[22]: ({'learning_rate': 0.5,
       'max_depth': 5,
       'min_child_weight': 2,
       'n_estimators': 20},
       0.619091448028456)

[23]: xgb_class: xgb.XGBClassifier = xgb_search.best_estimator_
xgb_class.fit(train_X, train_y)
xgb_pred = xgb_class.predict(test_X)
fig, ax = plt.subplots(figsize=(10, 10))
ConfusionMatrixDisplay.from_predictions(test_y, xgb_pred, ax=ax, display_labels=le.classes_)
accuracy_score(xgb_pred, test_y)

[23]: 0.62457798784605
```



Градиентный бустинг получился менее точным по сравнению с логистической регрессией, но он хотя бы выглядит реалистично (accuracy не 1).